

Signals and Terminal I/O

- Signals and how to control the program behavior in handling signals.
- Terminal I/O
- Readings
 - APUE 10.2-10.5, 10.9 - 10.14

Signals

- A form of inter-process communication
- tells a process that some event occurs
 - the kill command
 - “kill -l”
 - “kill -s INT pid”
 - “kill -INT pid”
 - Type “ctrl-C” when a program is running (SIGINT)
 - Memory violation (SIGSEGV), etc
 - Divided by 0 (SIGFPE).
 - A process dies (SIGHUP and SIGCHLD).
 - a packet arrives
 -

Some Commonly Used Signals

- SIGABRT, SIGALRM, SIGCHLD, SIGHUP, SIGINT, SIGUSR1, SIGUSR2, SIGTERM, SIGKILL, SIGSTOP, SIGSEGV, SIGILL
- All defined in `signal.h`
- `man -s 7 signal` on `linprog`

Signal

- When a process receives a signal, it performs one of the following three options
 - Ignore the signal
 - Two signals cannot be ignored
 - SIGKILL and SIGSTOP
 - Perform the default operation
 - Ignore the signal
 - exit
 - Catch the signal
 - Informs kernel to call user-defined function when signal occurs
- We can also block a signal from happening
 - Kernel remembers if a signal occurs and deliver it when we unblock the signal

Catch a Signal

- Similar to interrupt (software interrupt)
- When a process receives a signal:
 - stop execution
 - call the signal handler routine
 - continue
- Signal can be received at any point in the program.
- Most default signal handlers will exit the program.

ANSI C **signal** Function

- **syntax:**
 - #include <signal.h>
 - void (*signal(int signo, void (*func)(int)))(int);
- **semantic:**
 - signo -- signal number (defined in signal.h)
 - func: SIG_IGN, SIG_DFL or the address of a signal handler
 - SIG_IGN: ignore signal
 - SIG_DFL: perform default action
 - Handler may be erased after one invocation.
 - How to get continuous coverage?
 - Still have problems – may lose signals
- **See example1.c (lingprog and program): not well defined**

Block/unblock Signal: sigprocmask

- **Manipulate signal sets**

- #include <signal.h>

- int sigemptyset(sigset_t *set);

- int sigfillset(sigset_t *set);

- int sigaddset(sigset_t *set, int signo);

- int sigdelset(sigset_t *set, int signo);

- int sigismember(const sigset_t *set, int signo);

- **Manipulate signal mask of a process**

- int sigprocmask(int how, const sigset_t *set, sigset_t *oset);

- How: SIG_BLOCK, SIG_UNBLOCK, SIG_SETMASK

- See example2.c, example2a.c, example3.c

Critical Region

- For a critical region where you don't want certain signal to come, the program will look like:

```
sigprocmask(SIG_BLOCK, &newmask,  
            &oldmask);  
..... /* critical region */  
sigprocmask(SIG_SETMASK, &oldmask, NULL);
```


The `sigaction` Function

- Supersedes the `signal` function + signal blocking
- `#include <signal.h>`

```
int sigaction(int signo, const struct sigaction * act,  
struct sigaction *oact)
```

```
struct sigaction {  
    void (*sa_handler)(); /* signal handler */  
    sigset_t sa_mask; /*additional signal to be block  during  
                        execution of signal handler*/  
    int sa_flags;        /* various options for handling signal */  
};
```

- See `example4.c` and `example4a.c` (noted the program behavior with multiple signals)

The `kill` Function

- Send a signal to a process
- `#include <signal.h>`
- `#include <sys/types.h>`
- `int kill(pid_t pid, int signo);`
 - `pid > 0`, normal
 - `pid == 0`, all processes whose group ID is the current process's group ID.
 - `pid == -1`, all processes for which sender has permission to send
 - `pid < -1`, all processes whose group ID = `|pid|`
- See `example5.c`
- See `example6.c` for the use of `alarm`.

Impact of Signals on System Calls

- A system call may return prematurely
- See `example7.c`
- How to deal with this problem?
 - Check the return value of the system call and act accordingly
 - Check the `errno` variable.
 - See `example7a.c`

Terminal I/O

- The semantics of an output operation is relatively simple. Input is rather messy.
- Two input modes:
 - Canonical mode: the default mode, input line by line
 - Noncanonical mode: input characters are not assembled.
- We will focus on noncanonical mode
 - When do we use it?
 - What should we expect when using this mode for input?
 - Example: Which input mode does vi use?
 - How about an “old” shell such as Bourne Shell?
 - How about a new shell such as bash that supports command-line editing?

The `termios` Structure

- In POSIX.1, all the characteristics of a terminal device that we can examine and change are in a `termios` structure (`termios.h`)

```
struct termios {  
    tcflag_t c_iflag; /* input flag */  
    tcflag_t c_oflag; /* output flag */  
    tcflag_t c_cflag; /* control flags */  
    tcflag_t c_lflag; /* local flags */  
    cc_t     c_cc[NCCS]; /* control characters */  
}
```

Manipulating termios Structure

- Functions to get and set the fields in the termios structure
 - tcgetattr and tcsetattr;
 - #include <termios.h>
 - int tcgetattr(int fildes, struct termios *termios_p)
 - int tcsetattr(int fildes, int optional_actions, const struct termios *termios_p)
 - optional_actions: TCSANOW, TCSADRAIN, TCSAFLUSH

Noncanonical Mode

- Turn on the noncanonical mode:
 - Unset the ICANON flag in c_lflag
 - `myterm.c_lflag &= ~ICANON`
 - When will a read return using the noncanonical mode for input?
 - Number of characters (VMIN)
 - Time (VTIME)
 - Specified in the c_cc field
 - `c_cc[VMIN] = ???, c_cc[VTIME] = ???`
 - `VMIN > 0, VTIME > 0`
 - `VMIN = 0, VTIME > 0`
 - `VMIN > 0, VTIME = 0`
 - `VMIN = 0, VTIME = 0`
- See `example8.c` and `example9.c`