

Interprocess Communications

- System V interprocess communication (IPC) mechanism
 - Message Queues
 - Semaphores
 - Shared Memory
- Summary of processes related concepts
- Readings
 - APUE 15.6 —15.9

Why Other IPC Mechanisms

- Pipes/sockets
 - FIFO semantics
- Signals: sending flags
- Sometimes, we want something beyond FIFO
 - FIFO with tags (message queue)
 - File semantics: the content is always there unless it is modified explicitly (shared memory)
 - Once concurrency is allowed in shared data, we will need a way to protect (lock) the data (semaphore)

Message Queues

- What are they?
 - Similar to the FIFO pipes, except that a tag (type) is matched when reading/writing.
 - Allowing cutting in line (I am only interested in a particular type of message)
 - Equivalent to merging multiple FIFO pipes in one.

Message Queues

- **Creating a message queue:**
 - `int msgget(key_t key, int msgflag);`
 - Key can be any large number. But to avoid using conflicting keys in different programs, use `ftok()` (the key master).
 - `key_t ftok(const char *path, int id);`
 - Path point to a file that the process can stat
 - Id: project ID, only the last 8 bits are used
 - Msgflag controls how message queue should be created
 - Access mode
 - What if a queue exists with the same key

Message Queues

- A linked list of messages stored within the kernel and identified by a message queue identifier.
 - Every message has a type field, and a nonnegative length, and the actual data bytes.
 - `msgsnd` puts a message at the end of the queue
 - `msgrcv` gets a message, may not follow FIFO order (can be based on type)
 - Has resource limits: MSGMAX, MSGMNB, MSGMNI, MSGTQL.

Message Queues

- Message queue operations

```
int msgget(key_t, int flag)
```

```
int msgctl(int msgid, int cmd, struct msgid_ds *buf)  
    cmd: IPC_STAT, IPC_SET, IPC_RMID
```


```
int msgsnd(int msgid, const void *ptr, size nbytes, int flag);  
    flag: IPC_NOWAIT
```

```
int msgrcv(int msgid, void *ptr, size_t nbytes, long type, int flag);  
    type = 0: first msg in queue  
    type > 0: first msg whose msg type == type  
    type < 0: first msg whose msg type is lowest value <= |type|
```

- Performance advantage is no longer there in newer systems (compared with pipe)

Shared Memory

Common chunk of read/write memory among processes



Creating Shared Memory

```
int shmget(key_t key, size_t size, int shmflg);
```

Example:

```
key_t key;  
int shmid;
```

```
key = ftok("<somefile>", 'A');
```

```
shmid = shmget(key, 1024, 0644 | IPC_CREAT);
```

– See example `shm_create.c`

Attach and Detach Shared Memory

```
void *shmat(int shmid, void *shmaddr, int shmflg);  
int shmdt(void *shmaddr);
```

Example:

```
key_t key;  
int shmid;  
char *data;  
  
key = ftok("<somefile>", 'A');  
shmid = shmget(key, 1024, 0644);  
data = shmat(shmid, (void *)0, 0);  
...  
shmdt(data);
```

– See example shm_attach.c

Deleting Shared Memory

```
int shmctl(int shmid, int cmd, struct shmid_ds *buf);
```

```
shmctl(shmid, IPC_RMID, NULL);
```

- See example `shm_delete.c`

Semaphores

- Managing concurrent access to shared memory segment.
- Using Semaphores
 - Creation: `semget (...)`
 - Incr/Decr/Test-and-set : `semop (...)`
 - Deletion: `semctl(semid, 0, IPC_RMID, 0);`

Online tutorial

<http://beej.us/guide/bgipc/output/html/multipage/semaphores.html>

Command-line IPC control

- **ipcs**
 - Lists all IPC objects owned by the user
- **ipcrm**
 - Removes specific IPC object

Summary of Process Related Concepts

Process Environment

- Command line arguments: argc, argv
- environ and getenv()
- getpid(), getuid(), getppid()
- How does a program access the second (command line) argument?
- How does a process access the variable you set in shell using commands such as “setenv TERM vt100”?
- How does a process know its parent’s process id? How does a parent know its children’s process ids?

Process Management

- fork, exit, wait, waitpid, execv
- How can a parent process know whether its child has executed successfully?
- How to determine whether execv runs a command successfully?

File Operations

- What are the related data structures for file operations?
- open, close, read, write, unlink, dup.
- How to redirect the standard input/output/error?

Interprocess Communications

- Pipe
- What kind of processes can communicate with pipes?
- How to implement “ps | grep duan | more”?
- Message queue
- Shared memory
- Semaphore

IPC

- Signal
- What is the typical default action for a signal?
- Blocking/unblocking a signal
 - sigset manipulation
 - sigfillset, sigemptyset, sigaddset, sigdelset, sigismember
 - the sigprocmask system call
- Install signal handler (can ignore a signal or use default handler)
 - signal
 - sigaction
- Sending signal: kill, alarm

Terminal I/O

- canonical mode
- noncanonical mode
 - tcgetattr, tcsetattr
 - termios data structure

Process Group, Session, and Controlling Terminal

- Related to job control
 - Who gets to access the keyboard? Who to send signal generated from the keyboard.
- Foreground and background processes
- Joining a group or creating a group: `setpgid`
- Making a group foreground or background
 - `tcgetpgrp/tcsetpgrp`